



AI Product Delivery Transformation Playbook v2.0

Optum Technology

AI Tiger Team | September 2026

Table of Contents

- 1 Executive Overview
- 2 AI Product Delivery Lifecycle
- 3 New Team Model
- 4 Success Metrics
- 5 Technology Enablement
- 6 Transformation Onboarding

1

Executive Overview

The Source of Truth

What the AI Tiger Team is, what this playbook is, why it exists, and how Optum Technology is using it to drive a coordinated, enterprise-wide shift to AI-powered product delivery.

AI Tiger Team: Accelerating Optum Tech's AI Transformation

Driving adoption, workforce evolution, and new ways of working at scale.



Integrate AI Seamlessly to Amplify Value

Embedding AI into our daily workflows seamlessly so it feels like a superpower, enabling exceptional value with greater speed and ease.



Empower Our People to Thrive in an AI-Powered Future

Turning uncertainty into opportunity. Providing a clear vision for evolving roles and pathways to gain the skills needed to thrive and accelerate career growth.



Cultivate a Culture of Trust and Continuous Evolution

Leading with transparency, empathy and commitment to responsible AI use to guide teams through change with clear strategy and feedback.

35+ **AI Practitioners** – Optum Tech peers are leading the change

Supported by **Enablement Partners** from Applied AI, ETIPS, OTU, ESRO and more!

The AI Tiger Team Directly Supports the 2026 UHG AI Agenda

AI Tiger Team flagship deliverables are directly advancing the three UHG AI core approaches.

1 — OPTIMIZE TODAY'S WORK

- Replacing unstructured AI experimentation with a consistent, repeatable workflow every team follows
- AI agents at every phase reduce friction and cycle time across the highest-volume delivery work

- Eliminating handoff overhead and dependency delays — reducing rework and cycle time across the highest-volume work
- Proven: UHC Member Care — 70%+ productivity gain with 50% smaller teams

2 — BUILD REUSABLE FOUNDATIONS

- GitHub Spec-Kit delivers AI PDLC as an enterprise-approved plug-in for Claude Code, Codex, and Copilot
- Optum AI Catalog provides governed agents, skills, and prompts every pod inherits at each phase

- UAIS governance and the AI commons provide the policy guardrails and shared learning foundation every pod runs on
- CI/CD connectors (MCP, A2A) wire pods into enterprise systems without custom integration work

3 — REIMAGINE THE MODEL

- Fundamentally redesigning how software is conceived, built, and run — AI executes, humans own outcomes
- Permanently shifts cost and economics across 30,000 Optum Tech roles

- Replacing role-fragmented teams with 3–5 person pods with full end-to-end accountability — 10 specialists become 4 AI-augmented generalists
- Permanently resetting team economics and org structure across Optum Technology

The AI Transformation Story In One View

Transforming how 30,000 Optum Technology roles deliver product — from traditional SDLC to AI-powered delivery.

WHERE WE ARE TODAY

Traditional SDLC

- Teams experimenting with AI individually — no structure
- Tools arriving without workflow guidance
- No shared measurement of what's working
- Roles defined for a pre-AI world
- Large, fragmented teams with handoff overhead



WHERE WE'RE GOING

AI-Powered Delivery

Small, outcome-focused pods

Cross-functional, end-to-end ownership

AI drives the lifecycle

Humans set direction and own outcomes

Shared metrics

One standard for measuring what's working

Structured on-ramp

Tools, coaching, and accountability to scale

80% of teams in the new model within 12-18 months

5% today → 15% in 1-2 months → 20% in 4-6 months → 40% in 6-8 months → 60% in 8-12 months → 80% in 12-18 months

What is changing because of AI Product Delivery Transformation

This is not a technology change. It is a new operating model for how teams create value, organize work, and measure outcomes.

Optimize for flow, not utilization

- Success is how fast outcomes reach production — not how fully each person is booked.
- Capacity flexes to the work; people shift to where the outcome needs them.

Organize around outcomes, not capabilities

- Work is organized around a change in customer, user, or business behavior — not an application to build.
- We minimize effort that isn't contributing to that outcome, so the highest-value change moves first.

Pods are temporary, assembled from a shared pool

- People belong to a pool of skills first, then join a pod when an outcome needs them.
- A pod lasts only as long as its outcome; when it ships, people flow to the next priority.
- Members often report to different managers, and the pod directs its own work.

The pod owns the whole outcome while it exists

- One pod carries the outcome end to end across every stack it touches — for the life of that outcome.
- Dependencies are resolved inside the pod, so coordination overhead drops.

Measure the outcome, not the pod

- Success is the value delivered and everyone who contributed to it — across pods and managers.
- Velocity and per-team activity no longer signal progress; they measure motion, not value.

Let ownership follow the work, not the pod


- When a pod dissolves, its context is codified into reusable skills, so knowledge stays.
- Whoever built it still owns it wherever they move next, keeping accountability intact.

AI Transformation Accelerators Provide the Strongest Signal for Scale

Transformation Accelerators are most closely aligned to the New Team Model. By identifying them, along with divisional scaling plans, we can establish onboarding priorities and target teams best positioned to scale.

Earlier in Transformation Journey

Further Along in Transformation Journey




2

AI Product Delivery Lifecycle

The Model

What the new way of working looks like — the phases, the AI agents, and how work flows from opportunity to value realization.

AI Product Delivery Lifecycle



Phase Deep-Dive | Portfolio Discovery

Shaping and prioritizing portfolio opportunities and making go / no-go investment decisions

What Happens in this Phase

- Defining problem spaces and desired outcomes
- Prioritizing opportunities based on value and feasibility
- Sequencing work across the portfolio
- Making go / no-go investment decisions

Key Skills Required

- Outcome framing and value definition
- Tradeoff-driven prioritization
- Market and customer signal synthesis
- Portfolio planning and decision facilitation

AI Tools & Agents

- Market Insights and Opportunity Scout agents
- Feasibility analysis and risk identification
- AI-assisted prioritization scenarios
- Drafting opportunity definitions and metrics

Success Looks Like

- Only high-value opportunities move forward
- Clear, prioritized backlog aligned to strategy
- Investment decisions are fast and evidence-based
- No unprioritized work enters Inception

What Changes From Today

- Less intake, more active portfolio shaping
- Work defined by outcomes before funding
- Decisions centralized versus distributed ambiguity
- No pods formed/assigned until value is clear

Phase Deep-Dive | Inception

Defining complete solution scope and confirming readiness to build before construction begins

What Happens in this Phase

- Defining complete, unambiguous solution scope
- Translating outcomes into executable specifications
- Aligning cross-functional stakeholders
- Removing ambiguity before build begins

Key Skills Required

- Structured problem decomposition
- Solution architecture alignment
- Acceptance criteria and contract definition
- Cross-functional facilitation

AI Tools & Agents

- Spec Author and Solution Architect agents
- Rapid prototyping and scenario generation
- AI-assisted decomposition and acceptance criteria
- Gap and inconsistency detection

Success Looks Like

- Spec is complete, testable, and unambiguous
- No guessing during Construction
- Stakeholders aligned before build
- Work ready for immediate execution

What Changes From Today

- Less clarification during development
- More upfront clarity and alignment
- Specs owned and validated collaboratively
- Build does not start until definition is complete

Phase Deep-Dive | Construction

Building and continuously validating the solution into production-ready increments

What Happens in this Phase

- Delivering working, validated solution increments
- Continuous validation against requirements
- Maintaining alignment to outcome intent
- Producing production-ready outputs

Key Skills Required

- Rapid iteration and feedback loops
- Validation-first engineering mindset
- Code quality and architecture adherence
- Continuous integration and delivery

AI Tools & Agents

- Coding Copilot and Code Reviewer agents
- Automated test generation and execution
- Continuous validation and defect detection
- AI-assisted iteration and refinement loops

Success Looks Like

- Features validated as they are built
- No late-stage testing bottlenecks
- High confidence in solution correctness
- Output ready for deployment without rework

What Changes From Today

- Build and test happen simultaneously
- No dev-to-QA handoffs
- Frequent demos replace status reporting
- Validation drives progress

Phase Deep-Dive | Readiness

Proving the solution is safe, compliant, and operable before deployment

What Happens in this Phase

- Ensuring solution is safe and operable in production
- Validating compliance, security, and quality requirements
- Confirming operational readiness (monitoring, runbooks)
- Final go/no-go deployment decision

Key Skills Required

- Risk assessment and decision-making
- Operational readiness evaluation
- Compliance and security validation
- Deployment planning and rollback strategy

AI Tools & Agents

- Compliance Checker and Release Manager agents
- Telemetry and monitoring setup automation
- Readiness validation and gap detection
- Security and policy enforcement checks

Success Looks Like

- Deployment decisions are evidence-based
- No operational surprises post-release
- Monitoring and ownership fully defined
- Risks explicitly accepted or resolved

What Changes From Today

- No checklist-based deployments
- Deployment blocked without evidence
- Ops readiness built earlier in lifecycle
- Clear ownership before go-live

Phase Deep-Dive | Operations

Operating the system reliably and driving continuous improvement of outcomes

What Happens in this Phase

- Maintaining system stability and reliability
- Monitoring performance and usage
- Driving continuous improvement
- Feeding insights back into the lifecycle

Key Skills Required

- Incident management and root cause analysis
- Observability and telemetry interpretation
- Performance and cost optimization
- Outcome measurement and value analysis

AI Tools & Agents

- Production Monitor and Incident Responder agents
- Automated anomaly detection and RCA support
- Cost optimization and usage insights
- AI-driven trend and pattern identification


Success Looks Like

- System reliability meets or exceeds SLOs
- Issues detected and resolved quickly
- Measurable outcomes achieved in production
- Continuous flow of improvements

What Changes From Today

- Operations drives improvement, not just stability
- Production data drives decisions
- Continuous feedback into Discovery and Inception
- Value proven in production, not assumed

Outcome based work



Next iteration of this view will organize the flow by accountability / cadence.

3

New Team Model

The Who


How teams are structured around the work — pods, roles, accountability, and the shift from role-fragmented to outcome-oriented delivery.

AI Product Delivery — Value Stream Role Ecosystem

One model. Clear roles. Clear interactions. Shared outcomes.

INTERACTION KEY

- Assign
- > Influence
- ...> Feedback
- Business Value
- Delivery Environment
- Delivery Ownership
- Enablement
- Flow & Orchestration
- Ops Ownership



New Team Model: From traditional scrum teams to AI-enabled pods

The AI-enabled pod is the new standard delivery unit. Shifting from functional teams to AI-enabled, cross-functional pods that own outcomes end-to-end.

The New Team Model is a single, AI-powered pod of four roles, collectively accountable for delivery outcomes.

Product

Requirements · Prioritization · Prototype

Builder

Design · Code Quality · Security

Quality

Quality Strategy · Testing Quality ·
Release Gate Enforcement

Operations

SLOs · Incident Response ·
Observability

TODAY	
Scrum Team (~10-15 people)	
Team size	10-15 people (organizational formula)
Roles	Product Owner, Scrum Master, TPM, Tech Lead, Architect, Developers, QA, Operations (distinct roles)
Location	Globally distributed across multiple time zones
Ownership	Distributed across teams; outcomes split across multiple backlogs
Coordination	Cross-team alignment and negotiation; dependencies discovered late
Execution	Sequential and fragmented by 2-week sprints; staggered releases across teams
Tools	Assisted with basic AI Productivity Tools



TOMORROW	
Assembled Pod (~3-5 people)	
Team size	3-5 people (pod size adapts to work)
Roles	Product, Builder, Quality, Operations (pod member can have multiple roles)
Location	Pod is geo-located (+ - 3 times zones of each other)
Ownership	Unified within the pod; one outcome built end-to-end
Coordination	Happens naturally inside the pod; dependencies resolved up front
Execution	Parallel, continuous and integrated; built and deployed together
Tools	Agentic AI toolkit with standardized skills repository augmented with digital workers

The Optum AI Product Delivery Pod: Our Standard Team Unit

Default pod size is 3-5 people with all necessary skills based on needs of the outcome — a small, accountable team, accelerated and powered by AI



Product

Accountable for:

- Confirming user needs and assumptions
- Prioritizing opportunities within the problem space
- Creating testable prototype or solution concepts ready for validation

Example Agents:

- Feature decomposition agent
- Spec Author/PRD Agent
- Feasibility Analyst



Builder

Accountable for:

- Technical design and code quality
- CI/CD pipeline
- Security and compliance

Example Agents:

- Reviewer agents
- Spec and design reviews
- CI/CD Triage Agent



Quality

Accountable for:

- Quality strategy across the lifecycle
- Shift-left: quality from story to ship
- Release gate definition and enforcement

Example Agents:

- Test Generator
- Coverage gap agents
- Edge case generator
- Acceptance Criteria Reviewer



Operations

Accountable for:

- SLO/SLI ownership and error budgets and feedback loops
- Incident response and post-mortems
- Observability and cost efficiency

Example Agents:

- Production Monitor
- Incident Responder
- Cost Optimizer
- Telemetry agent

1

End-to-end. Cross-functional. Complete.

Every tech and business role and skill needed to ship. All in one pod.

2

Pod size follows the work, not a formula.

One person can wear multiple hats — or share one hat with others. Smaller teams than today, but complexity of work determines size and number of pods.

Non-Pod Roles: Work Outside the Pods

Where each role originates relative to today, and what's different in the new operating model

Role	What we see today	What's Different in the New Model
Initiative Owner	Often seen as Portfolio Manager, Product Manager, or Capability Manager — owns the business roadmap	Works across the value stream, not one product/application; end-to-end outcome definition, AI-enabled processes
Delivery Lead	Often seen as Engineering Manager, Tech Lead; usually has people-management responsibilities; accountable for the delivery of the team	Not accountable for the work being delivered; accountable for the systems that enable delivery — e.g., ensuring the right expertise is codified into skills available to pods
Orchestration Lead	Split across Technical Program Management, Release Train Engineer (RTE), Scrum Master, Technical Product Management	Flow management is the central focus — how work moves through the system; may be split across multiple people in large orgs; accountable for capacity management of pods and pool
Enablement Layer	Embedded specialists (architects, UX, QE, DevOps) inside individual teams or supporting multiple teams	Focused on delivering agentic skills used across pods, i.e. if pods need better architecture, Enablement is accountable for improving that skill in the agentic flow, not for doing or approving the design
Shared Operations Pool	Work pulled into scrum teams between feature work, or handled by dedicated O&M/support teams	Outcome pods no longer absorb small fixes between features — that work now belongs entirely to the unassigned pool. There should always be some pod-capable people unassigned.

1

Roles, not headcount.

These are accountabilities, not job titles or positions. One person can hold several of these roles, and a single role can be shared across several people.

2

Every skill is still needed. This is work outside the pod.

The full set of tech and business roles and skills is still required to ship. These roles cover what happens around and between pods, not inside them.

Pod Role Deep-Dive | Product

Single-threaded owner of the problem space, outcomes, and delivery sequencing

Accountable for

- Customer and user problem ownership
- Business outcome accountability (outcomes, not output)
- Priority and trade-off decisions — What and When
- Stakeholder alignment across the lifecycle

Key Skills Required

- Continuous discovery: research, framing, validation
- Outcome-based road-mapping (JTBD, OKR alignment)
- Backlog ownership: stories, acceptance criteria, refinement
- Product metrics: instrumentation and funnel analysis

AI Tools & Agents

- PRD drafting, story generation, roadmap summaries
- Opportunity Scout and Market Insights agents
- Feasibility Analyst agent for rapid scoping
- AI-assisted planning and story breakdown

Success Looks Like

- Team ships measurable outcomes, not just features
- Discovery runs continuously — next bet always validated
- Stakeholders do not need a status meeting

What Changes From Today

- TPM and Scrum Master coordination absorbed into Product
- Less spec writing; more discovery and deciding
- Outcome ownership is single-threaded — one person

Pod Role Deep-Dive | **Builder**

Owens technical design, code quality, delivery pipeline, and production deployment

Accountable for

- Technical architecture and approach decisions
- Code quality, velocity, and engineering standards
- Delivery pipeline: CI/CD, environments, deployment
- Security and compliance

Key Skills Required

- Software engineering: design, implementation, review
- AI-augmented development: copilot-first workflows
- System design: APIs, data modeling, service boundaries
- DevOps: CI/CD, containerization, IaC fundamentals

AI Tools & Agents

- Agentic code completion and refactoring
- Task decomposition and code generation
- Code Reviewer and Coding Copilot agents
- Spec interpretation, design and PR review

Success Looks Like

- Higher throughput — AI handles first drafts and boilerplate
- Code reviewed and tested before humans touch it
- Deployment is a non-event — pipeline is fully automated

What Changes From Today

- SWE, DevOps, and design skills converge into one role
- Dev-to-ops handoffs eliminated within the pod
- AI copilot is the default starting point, not optional

Pod Role Deep-Dive | Quality

Embeds quality across the full lifecycle — from story definition through production

Accountable for

- Quality strategy: test approach, coverage, and risk prioritization
- Shift-left execution: quality at story, design, and build stages
- Release gate definition: criteria that make something shippable
- Continuous improvement of test automation and coverage

Key Skills Required

- Test strategy: unit, integration, contract, E2E, exploratory
- Automation engineering: frameworks and CI integration
- Risk-based testing: where coverage matters most
- Acceptance criteria review and story-level quality partnership

AI Tools & Agents

- Test Generator agent: test plans and cases from specs
- Coverage gap agents: undertested paths identified automatically
- Automated regression and smoke suites wired into CI/CD
- Edge case generation and acceptance criteria checks

Success Looks Like

- Defects caught at story or build stage, not in production
- Test coverage expands without proportional manual effort
- Release gate is automated, consistent, and trusted by the pod

What Changes From Today

- QA moves from end-of-sprint gate to lifecycle partner
- Manual test writing replaced by AI-generated, human-reviewed suites
- Quality owns release readiness, not just pass/fail

Pod Role Deep-Dive | Operations

Owens production reliability through SLOs, error budgets, and observability

Accountable for

- SLO and SLI definition — sets what "reliable enough" means
- Error budget management — enforces reliability/velocity trade-off
- Incident response and post-mortem process end-to-end
- Observability, telemetry, and production cost efficiency

Key Skills Required

- SLO/SLI design and error budget tracking (Google SRE model)
- Observability: metrics, logs, traces, and alerting design
- Incident management: on-call practice, runbooks, post-mortems
- Infrastructure as code and cloud platform operations

AI Tools & Agents

- Production Monitor agent: health and anomaly detection
- Incident Responder agent: AI-assisted triage and runbooks
- Cost Optimizer agent: spend analysis and right-sizing
- Telemetry Setup and Compliance Checker agents

Success Looks Like

- Error budgets tracked and used to make velocity decisions
- Incidents detected and triaged before user impact (MTTR declining)
- DORA stability metrics (CFR, TTR) in elite range consistently

What Changes From Today

- Ops shifts from reactive firefighting to proactive reliability
- SLOs replace "keep it up" as the definition of success
- On-call burden reduced via AI-first triage and auto-remediation

Role Deep-Dive | Initiative Owner

Business accountability owner for an initiative and its outcomes

Accountable for

- Define objectives, outcomes & success criteria
- Prioritize outcomes & approve investment
- Provide ongoing feedback throughout delivery
- Determine customer-readiness & closure timing

Key Skills Required

- Business domain expertise and value definition
- Portfolio prioritization & investment trade-offs
- Stakeholder communication & feedback loops
- Business case & value-readiness assessment

Key Interactions

- Sets value & success criteria for the Outcome Pod
- Receives progress & readiness signals from the Outcome Pod

Success Looks Like

- Outcomes measurably move the business forward
- Pods know what “done” means from day one
- Consistent, iterative feedback

What Changes From Today

- Business partners define outcomes, not solutions or feature backlogs
- Portfolio Discovery becomes the front door for demand
- Less component management, more value generation

Role Deep-Dive | Delivery Lead

Owner of the delivery environment and standards pods build within

Accountable for

- Sets engineering, delivery, quality & AI standards
- Defines practices & evolves the delivery environment
- Owns workforce readiness & capability development
- Drives reusable assets, skills & agent adoption

Key Skills Required

- Engineering & delivery standards-setting
- Technical coaching & capability development
- AI-augmented delivery practices (skills & agent adoption)
- Delivery environment & tooling strategy

Key Interactions

- Sets standards & guardrails for the Outcome Pod
- Receives feedback & improvement needs from pods
- Shares improved practices with the Shared Operations Pool

Success Looks Like

- Pods meet the quality bar without being told how
- New contributors ramp quickly on shared standards
- Reusable assets get adopted, not reinvented

What Changes From Today

- Means different things today — this model standardizes it
- Scrum-team reporting lines stay open to the transformation
- No people-management or work-assignment accountability
- Owns continuous improvement of the ecosystem, not pod outcomes

Role Deep-Dive | **Orchestration Lead**

Manager of how work and people flow through the delivery system

Accountable for

- Forms & staffs outcome pods
- Balances capacity & coordinates dependencies
- Routes escalations & optimizes flow
- Governs the Shared Operations Pool

Key Skills Required

- Capacity planning & staffing allocation
- Dependency mapping & flow management
- Escalation handling & conflict resolution
- Cross-pod coordination

Key Interactions

- Provides talent & capacity alignment to the Outcome Pod
- Receives requests for talent & flow support from pods
- Governs transitions of work & people with the Shared Operations Pool

Success Looks Like

- Pods rarely sit blocked on people or dependencies
- Staffing changes land without breaking momentum
- Bottlenecks don't surprise anyone — capacity is visible

What Changes From Today

- Consolidates Program Mgmt, RTE, Scrum Master & Technical PM
- Enables teams to work without dependencies, not manage them
- Owns flow & performance metrics for the whole value stream
- Defines pod assembly/reassembly — no people-management role

Role Deep-Dive | Enablement Layer

Reusable expertise shared across outcomes: architects, UX, ops/agent owners

Accountable for

- Maintains reusable standards & best practices
- Provides architecture & quality guidance
- Builds reusable skills, agents & prompts
- Curates reusable patterns

Key Skills Required

- Reusable-asset design (skills, agents, prompts)
- Architecture & quality advisory
- Coaching & influence without direct authority
- Pattern curation & knowledge management

Key Interactions

- Advises & shares reusable assets with the Outcome Pod
- Receives needs, feedback & lessons learned from pods
- Contributes reusable assets & expertise to the Shared Operations Pool

Success Looks Like

- Pods reuse existing skills & patterns, not from scratch
- Guidance gets sought out proactively, not imposed
- Delivery speeds up across many pods, not just one

What Changes From Today

- Specialists shift from executing to enabling delivery
- Expertise is codified, not held by a few people
- No approval authority — pods decide, enablement advises

Role Deep-Dive | Shared Operations Pool

Capacity buffer and rotational pool between outcome assignments

Accountable for

- Provides post-closure support & maintenance
- Delivers small enhancements & reduces tech debt
- Triage vulnerabilities for closed outcomes
- Buffers capacity between outcome assignments

Key Skills Required

- Broad, generalist engineering & operational skills
- Rapid context-switching across codebases
- Triage & prioritization under ambiguity
- Onboarding & mentoring less experienced contributors

Key Interactions

- Receives improved practices from the Delivery Lead
- Receives reusable assets from the Enablement Layer
- Staffed via transitions from the Orchestration Lead
- Takes on outcomes transitioned by the Outcome Pod

Success Looks Like

- Closed outcomes keep running without a dedicated pod
- Technical debt gets paid down between assignments
- People move between assignments without idle time

What Changes From Today

- Moves away from dedicated O&M teams toward more automated operations
- New formal capacity-management mechanism between assignments
- People rotate between outcome pods and the Shared Operations Pool

Role Deep-Dive | Outcome Pod

Cross-functional unit accountable for delivering an outcome end-to-end

Accountable for

- Decomposes features & designs the solution
- Implements, tests & deploys
- Provides operational support while active
- Measures, validates & owns outcome delivery

Key Skills Required

- End-to-end delivery across Product, Builder, Quality & Ops
- Cross-functional collaboration within one accountable unit
- Outcome measurement & validation
- AI-augmented delivery across the lifecycle

Key Interactions

- Aligns with the Initiative Owner on success criteria
- Adopts standards & assets from Delivery Lead & Enablement
- Gets talent & capacity alignment from Orchestration Lead
- Transitions to the Shared Operations Pool after closure

Success Looks Like

- The outcome verifiably solves its chartered problem
- The pod rarely escalates — decisions happen inside it
- Operational issues get fixed without a separate handoff

What Changes From Today

- Team: 10-15 people, distinct roles → 3-5 people, one pod
- Ownership: split backlogs → unified, end-to-end
- Execution: staggered releases → continuous, integrated

Role Deep-Dive | AI Transformation Lead

Embedded long-term owner of sustained local adoption and execution

Accountable for

- Local transformation roadmap, priorities, sustainment plan
- Leadership alignment & adoption outcomes
- Transformation maturity & measurement
- Escalation of barriers to progress
- Long-term transformation success

Key Skills Required

- Cross-boundary coordination and stakeholder alignment
- Persistent follow-through and friction removal
- Progress measurement and clear leadership communication
- Local credibility and influence without formal authority

Core Activities

- Takes handoff from the AI Tiger Team Navigator
- Align leaders, product teams, development teams, & partners
- Drive action on adoption barriers
- Reinforce new ways of working over time

Success Looks Like

- Adoption continues after onboarding concludes
- Maturity improves against agreed measures
- Barriers get addressed, mitigated, or escalated
- AI Product Delivery becomes normal operations

What Changes From Today

- Ownership shifts from temporary enablement to local accountability
- Transformation plans become actively managed roadmaps
- Adoption sustains without routine Navigator involvement

Role Deep-Dive | AI Tiger Team Navigator

Temporary catalyst who creates readiness and momentum, then hands off locally

Accountable for

- Readiness assessment & structured onboarding facilitation
- Leadership & business-track alignment
- Early adoption coaching & risk visibility
- Handoff preparation to local ownership

Key Skills Required

- Cross-functional facilitation and change enablement
- Rapid organizational context and trust building
- Risk, dependency, and barrier identification
- Turning ambiguity into practical transformation plans

Key Interactions

- Guide team, leadership, and business onboarding conversations
- Facilitate outcome roadmaps, dependency plans, and transformation plans
- Prepare the embedded AI Transformation Lead and coordinate knowledge transfer

Success Looks Like

- Readiness gaps and dependencies addressed
- Leaders and teams aligned on the transformation path
- Teams operate in the AI Product Delivery model
- Clean handoff to a prepared local lead

What Changes From Today

- Onboarding becomes a transformation journey, not just training
- The Navigator enables and facilitates, but does not own it
- Engagement ends with a deliberate handoff to local ownership

4

Success Metrics

How We Know It's Working

The shared metrics every pod runs on that measure throughput, stability, and continuous improvement.

Directional Success Metrics: How we will measure adoption and outcome impact

We will measure AI Product Delivery Transformation through two layers: (1) transformation adoption and (2) outcome impact through DORA and FCT. FCT/TTB provides an early indicator of productivity gains, while DORA metrics become more meaningful as adoption scales.

Metrics as of 7/8/26 (DORA: 6-month average, FCT: YTD)

Type	Metric	What It Measures	OT	CE – Sutha's Org*	% Improvement
DORA	Throughput	Deployment Frequency	10.7 days	8.2 days	23.4%
DORA	Stability	Change Failure Rate	3%	3%	0%
DORA	Stability	Failed Deployment Recovery Time	unavailable	unavailable	unavailable
DORA	Throughput	Lead Time for Changes	7 days	4 days	42.9%
FCT	Throughput	FCT / TTB (Time to Build)	6.9 weeks	5.0 weeks	27.5%

* 160 out of 1,200 team members are in 30 AI PDLC pods today, scaling to 600 by Q3, and 1,200 by EOY.

CE: Example of Transformation Accelerator

UHC Member Care Pilot

The Member Care Hub helps UHC M&R members close STARs care gaps and improve preventative care engagement.

A unified member profile API eliminated data silos by creating a single real-time member view for care teams and digital experience.

Using the AI PDLC approach, AI was embedded across the delivery lifecycle—from requirements and design through code generation, testing and deployment automation.

95%+

AI-Generated Development

Code, specifications, and test automation

70%+

Productivity Gain

Compared to baseline delivery expectations

50%

Smaller Teams

1 pod delivered vs. 2 pods planned

5


Technology Enablement

The Framework


The enterprise tools, agents, and shared infrastructure that every pod runs on — from agent harnesses and skill libraries to the GitHub Spec-Kit and UAIS governance layer.

AI Product Delivery Lifecycle Building Blocks

Humans set goals and guardrails · **prompts** directs **agents** · **skills** and **plugins** help deliver the work.




Tools that Enable Developer Services and AIDLC



How the United AIDLC framework fits together

One governed foundation, extended by every business unit, and consumed everywhere engineers and business users work



One governed foundation.

The Optum AI catalog is the enterprise source of truth for agents, skills, and connectors, governed under UAIS.

A common baseline.

Every team runs the same five AIDLC phases on the same Spec Kit cycle, so delivery is consistent enterprise-wide.

Built for every business unit.

Business units extend the baseline for their own needs instead of forking it, so improvements keep flowing to everyone.

Works where teams already are.

The same skills and plugins install into coding-agent harnesses and browser-based chat agents alike.

Improves by design.

Engineering teams share learnings back into the catalog through an inner-source loop, so every pod inherits the gains.

6

Transformation Onboarding

When & How

The sequencing — from today's pilots to org-wide adoption, with clear milestones, success criteria, critical enablers, and the path to scale.

AI Product Delivery Transformation Begins with Clear Team Expectations

Three pillars determine onboarding readiness, transformation commitment and ongoing accountability.

1

PILLAR 1

Readiness Prerequisites

- Designated AI transformation lead
- Leadership sponsorship secured
- Dedicated capacity to learn adoption
- Data & compliance readiness cleared
- Baseline AI knowledge & fluency
- Documented use cases
- Pre-transformation team assessment and sentiment baseline complete

Hard gate — all 7 required

2

PILLAR 2

Transformation Commitments

- Adopt AI Product Delivery Lifecycle
- Implement new team model
- Share-back to the AI commons
- Human review of AI-generated outputs
- AI retrospectives
- AI incident reporting

Signed Ways of Working Pledge

3

PILLAR 3

Measurement Obligations

- Track and report the common throughput and stability metrics
- Monitor AI token utilization
- Complete post-transformation team assessment and sentiment survey
- Measure and report team capacity trends

Monthly Reporting Required

AIDLC Team Training/Onboarding Journey

Preparing teams to begin using AIDLC in day-to-day delivery work.

1 STEP 1

Foundational Training

Build the foundation before attending bootcamp.

Team Model & Roles

- Pod model
- Product, Builder, Quality & Operations

Outcomes & Lifecycle

- Outcomes vs. features
- Lifecycle overview
- Work flow expectations

AI Components & Readiness

- Tools & access
- Agents, skills & assets
- Environment validation

Outcome: Shared understanding of AIDLC + validated access to required tools.

2 STEP 2

Lifecycle Simulation & Workshop

Experience a guided AIDLC example, then apply AIDLC to real work.

Guided Simulation

- Walk through a sample outcome
- Observe the end-to-end workflow
- See AI and human collaboration
- Understand key artifacts and decisions

Discovery - Readiness

- Define problem
- Create and refine specifications
- Multiple build & validation cycles
- Validate quality and release readiness

Outcome: Completed a guided AIDLC workflow using a team-specific use case.

3 STEP 3

Onboarding Closeout

Align on expectations and next steps.

Lessons Learned

- Key takeaways
- Remaining questions

Ways of Working Pledge

- Responsible AI use
- Human accountability
- Team commitments

Retrospective

- Feedback on onboarding experience
- Improvement opportunities

Outcome: Team is prepared to begin using AIDLC in everyday delivery work.

- ✓ Completed foundational training
- ✓ Validated tools and access

- ✓ Experienced AIDLC end-to-end
- ✓ Applied a real use case

- ✓ Established team ways of working
- ✓ Ready to begin using AIDLC

Manager's Path to AI Product Delivery Transformation

Ready

Week 1-2

1

This is where you do the most work.

- Meet your assigned Tiger Team Navigator
- Run the assessment and sentiment baseline across your full group
- Carve out a dedicated pod that meets the readiness prerequisites
- Protect that pod's capacity from competing demands
- Keep working the rest of your group to line up the next pod

Train

Week 3

2

Step back. Let the pod do the work.

- Your pod attends the multi-day bootcamp
- You make sure participants complete training prerequisites in advance
- You hold the line on dedicated capacity during the workshop
- You attend the closing ceremony when the pod signs the pledge

Coach

Week 4

3

Stay visible. Stay out of the way.

- Tiger Team Navigator runs check-ins with your pod at week 4.
- You review the pod's monthly metrics and remove organizational blockers
- You shield the pod from scope creep and competing priorities
- You start the Ready stage for your next pod in parallel

Sustain

Week 6

4

Your pod becomes your proof point.

- Your pod attends a bootcamp refresher.
- Pod operates independently with async coach support
- You use the pod's results to recruit and prepare the next wave
- You sponsor pod members as references for other teams in your group

Your involvement

What You Get from Us

Tiger Team Navigator
Assessment instrument
Sentiment survey
Pod selection criteria
Readiness prerequisites

Workshop facilitation
Training prerequisites for participants
Tools and access provisioning
Ways of Working pledge

Scheduled check-ins
Measurement dashboard
Monthly reporting templates
Blocker escalation path

Reference playbook
Async coach access

What We Need from You

Sponsorship and visible commitment
Time to complete the assessment
Honest pod selection against the criteria
Dedicated capacity carved out and protected


Participants prepped and present for all three days
Capacity held during the workshop
Real work staged for the pod to walk through
Presence at the closing pledge

Monthly metric reviews on your calendar
Active removal of organizational blockers
Scope discipline so the pod is not pulled off
Ready stage started for your next pod

Pod treated as the standard, not the exception
Recognition for pod members as internal references
Next wave of pods queued from your group

The AI Tiger Team Website is Live

Come explore your one-stop shop for AI transformation resources, role-based training, and ways to get involved.



New Tiger Team site: Want to get involved or learn more? Our new site is LIVE! aitigerteam.optum.com

Role-based training: Can be found at otu.optum.com and on **MyLearning**

More AI resources: and visit the **AI Hub** on Sparq, **AI Dojo** at OTU, and **AI Guild** on Viva Engage

Appendix

SDLC → AI-Assisted SDLC → AI PDLC



HUMAN ROLE: Every task, every review, every approval — humans execute and inspect each step across every phase

AI ROLE: Minimal. No meaningful delegation of execution or decision.

HUMAN ROLE: Reviews and approves AI output, makes quality and risk calls, owns outcomes and accountability

AI ROLE: Drafts specs, generates code & tests, runs scans, assembles compliance evidence — accelerates every phase

HUMAN ROLE: Sets intent & outcome criteria, defines policies & guardrails, reviews outcome summaries, governs exceptions — architect of the system, not the executor

AI ROLE: Executes the full delivery lifecycle — building, testing, compliance validation, documentation, deployment, and monitoring — continuously and autonomously under human-defined guardrails